Development of a 3D display

Balduin Dettling

August 20, 2016

Preface

Motivation

For as long as I can remember, I've been fascinated by all kinds of technology. It was clear to me that my final thesis would be a practical one in the field of engineering. However, I didn't have anything specific in mind for quite some time.

One day, I stumbled upon a project report on the internet about a so-called persistence of vision display. It was a simple row of LEDs, driven by a microcontroller, and mounted on an electric motor. The contraption was (and presumably still is) able to draw images into the air by precisely controlling the row of LEDs. Fascinated by this, I did some research and found that a plethora of similar projects exists, some even with detailed instructions on how to build the device in question.

I had the idea of building several rows of LEDs and arranging them on top of each other. That way, I could display several two-dimensional bitmap images above another, resulting in a three-dimensional image.

Having had no previsous experience in the field of electronics, it seemed a bit daunting to build a 3D display. I could, however, motivate myself to attempt it anyway by telling myself that the project wouldn't be much more complex than all of the 2D versions out there (which turned out to be not very far from the truth).

Credits

First and foremost, I would like to express my gratitude to Mr. Patrick Spengler, my physics teacher and the mentor of this project. Not only did he solve numerous malfunctions, but he also is partly responsible for my interest in electronics and engineering.

I am thankful to Ms. Ursula Schamberger, teacher of the woodworking class, where I built the wooden frame. Likewise, I would like to thank Mr. Hanspeter Rieder, who kindly granted me access to the school's electronics workshop.

Last, but most certainly not least, I am indebted to Prof. Heinz Domeisen. He mentored the advancement and refinement of the thesis in preparation for the national competition organised by "Swiss Youth in Science".

Declaration of authorship

I hereby declare that the following thesis has been authored entirely by myself, Balduin Dettling, and confirm that all external sources of information are cited correctly.

Location and Date   Signature

Contents

1. Introduction

1.1 Current State of The Art

As discussed before, there is no shortage of projects that are similar to mine, but only display 2D images [1, 2, 3]. The name "POV display", short for persistence of vision display, has become commonplace. Another popular name is "Propeller Clock" — Propeller because it rotates like one, and clock because the round shape is ideal for displaying an analog clock.

As far as I know, however, only two people have built a working version of a three-dimensional POV display so far and have published it online [4, 5].[1] Neither of these is being produced on a grand scale, instead they are — like most 2D versions — side projects of hobbyists or engineers.

1.2 Goals

The objective of this project is to build a working prototype of a 3D persistence of vision display and to document the process of doing so. The necessary steps are (roughly) as follows:

- Understand and explain how and why such a display can work in principle

- Come up with realistic technical specifications, founded on the previously outlined theory

- Select parts and build the device according to the specifications

- Write a program that makes the device display a 3D image

- Draw conclusions: Did I achieve my goals? What could have been done better? How can the project be further improved?

---

[1]It turned out that there was a third one - When I wrote a forum post about the display, somebody made a comment about a similar display he built years ago: `https://www.reddit.com/comments/2q9sg6/-/cn4i4rk`

2. Theory and Planning

2.1 Principle of Operation

2.1.1 Persistence of Vision

Persistence of Vision refers to the optical illusion that results whenever the image we look at quickly changes. In such situations, we don't immediately see the new image, instead we continue seeing the old image for a short timespan after it has vanished. If wikipedia is to be believed, this timespan is about $\frac{1}{25}$ of a second, but of course it is going to be different for each pair of eyes [6].

Although we don't normally see this effect in action, there are situations where it becomes apparent. For example, if a car drives by at nighttime, we see traces of light behind the actual car lights for a short moment (Assuming our eyes don't follow the car and stay still relative to the environment). In general, bright flashes of light followed by (relative) darkness tend to be the most pronounced manifestation of persistence of vision. It is exactly this situation that can be recreated and exploited for the successful operation of a persistence of vision display.

2.1.2 Exploiting Persistence of Vision

Normally, Persistence of Vision is nothing more than an interesting effect. However, we can use this weakness of the human eye to our benefit.

If a light source is spun quickly enough in a circle, we don't see a spinning point anymore, but instead we see one solid circle. This is because the rotational period is shorter than the time it takes for the light to fade away on our retinas. The rotational period at which we stop seeing a flickering circle and start seeing a solid, still circle is called the flicker fusion threshold.

By controlling this light source quickly and accurately, we can make it so it is always turned on at certain points in the circle, and always turned off at others. By doing so, we are effectively multiplexing a zero-dimensional display (a point-like light source, e.g. an LED) so that it displays a one-dimensional image. By rotating a whole row of LEDs (a one-dimensional display), we can display two-dimensional images. Finally, if there are multiple rows of LEDs that together form a two-dimensional display, it is possible to display a three-dimensional image.

2.2 Specifications

- Resolution of $16 \cdot 100 \cdot 10$ pixels (radius $\cdot$ circumference $\cdot$ height)

- Each pixel consists of an RGB LED

- Colour depth of 3 bit per pixel, or one bit per LED

- Rotational frequency of at least 30 Hz, or as fast as it needs to be in order to not flicker

2.3 Components

2.3.1 Control

For controlling the LEDs, I have bought a Teensy 3.1 microcontroller development board [7]. It has a small footprint — ideal for the fast rotation — and still packs quite some power. The processor runs at 96 MHz at a register width of 32 bit, and there are 34 configurable I/O pins.

The controller has to have some kind of reference point, so that it can calculate the current speed and angular position of the device. For this, I used a simple hall sensor on the rotating side, which passes by a stationary magnet after each rotation.

2.3.2 LED drivers

34 I/O pins are quite a lot, but not quite enough to directly control 480 LEDs. In order to still access all those LEDs individually, I had to use LED drivers. These come in the form of integrated circuits and receive some amount of data that they use to control a corresponding amount of LEDs.

Initially, I had planned to implement a colour depth of 8 bits, allowing for $2^{24} = 16777216$ different colours. However, there were several difficulties that I would have had to deal with. First of all, all of the available LED drivers had one or both of the following shortcomings:

- Controlled using I^{2}C instead of SPI

  Because I^{2}C uses pullup resistors instead of a push-pull drive, it takes a little moment for the logic level to return to its normal state. In order to ensure correct data transmission, it is necessary to set a relatively slow clock speed. Even with an optimised library [8], the two I^{2}C interfaces of the Teensy are limited to 2.4 MHz if the clock speed it kept at 96 MHz — and because I^{2}C requires some overhead data to send out the slave's addresses, the actual data rate will be even lower.

  However, the minimum data rate required to control the LEDs at an 8-bit colour depth is as follows:

  $$24 \tfrac{\text{bit}}{\text{pixel}} \cdot 16000 \tfrac{\text{pixels}}{\text{rotation}} \cdot 30 \tfrac{\text{rotations}}{\text{second}} = 11.52 \,\text{Mbps}$$

  This means that I'd need at least 5 parallel I^{2}C lines to control all the LEDs with 8 bit colour depth. Because each Teensy only has two of them, I'd need three Teensy boards, which would be possible but very cumbersome. The SPI lines, however, can send a bit at every other clock cycle, leading to a bit rate of 48 MHz.

- Insufficient PWM frequency

  If the colour depth is more than one bit, the LEDs will have to be dimmed using PWM. Normally, a rather low frequency of a few hundred to a few thousand Hz is sufficient as a PWM frequency, and most LED drivers modulate at frequencies within that range. In my case, however, the LEDs are moving into a new pixel 3000 times a second (assuming rotationat 30 Hz). This means that a PWM frequency under 3000 Hz will involve losing information, and one slightly above 3000 Hz still won't look good. There would have to be a significant number of PWM cycles within each pixel, for example 10, which would require a PWM frequency of 30 KHz.

Because of these problems, I ditched the idea of 8-bit colour depth and instead went for the simpler approach of one-bit colour depth. Both problems are solved by switching to 1 bit: The data rate doesn't have to be as high (and there are lots of chips that are controlled by SPI), and PWM is not necessary anymore.

After quite some searching around, I decided to use the TLC5927 by Texas Instruments [9, 10]. It is essentially a shift register with some extra features. Like every normal shift register, it has a serial data input (SDI), a serial data output (SDO), and a clock pin (CLK). Additionally, there is a latch pin (LE), which serves the purpose of only refreshing the state of the LEDs once all data has been sent.

Furthermore, there is a pin called "output enable" (OE). If it is at GND, all LEDs are turned off, and it's at V_{DD}, they are controlled by the data that the driver chip has received.

The most important feature, however, is the fact that each output is controlled by a constant current driver. Without these, it would've been necessary to solder a resistor in series with each individual LED, for a total of 480 resistors, 960 additional solder joints and a whole lot of wasted board space. With the constant current drivers, however, I could use a single resistor to set the current for all 16 LEDs attached to one chip.

Figure 1: Pin layout of the TLC5927 LED driver

Unlike most simple shift registers, the drivers' outputs are current sinks instead of current sources. This will be important when selecting the LEDs: In order to control each colour separately in an RGB LED, the cathodes have to be separated, while the anodes can be connected together.
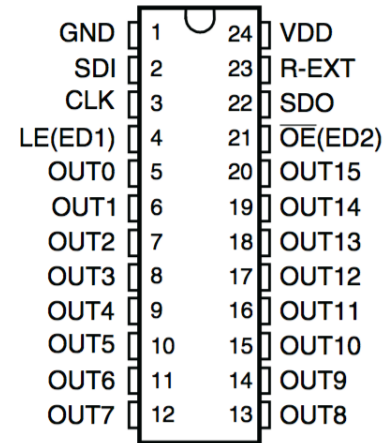
2.3.3 LEDs

Like all the other parts, I ordered the LEDs from Digikey. Their website contains an excellent tool, enabling the user to restrict the selection based on different criteria. My requirements were the following:

- Colours: Red, green and blue

- Current: 20 mA or more. Because the LEDs are within one pixel only $\frac{1}{100}$ of the time, they are essentially pulse width modulated, reducing their effective brightness. In order to still see the image clearly without darkening the room, they have to be bright enough.

- Mounting type: Surface mount, so I can still use the other side of the board

- Size: At most 5 mm by 5 mm, preferably smaller, but still large enough to solder by hand

- Diffused lens, so the colours mix well even when viewing the LED directly instead of using it for illumination.

These requirements reduced the category "LED Indication - Discrete" from 18940 to 130 items. From those remaining products, I chose the one that was cheapest in a quantity of 160. This led me to an LED with the beautiful name CLVBA-FKA-CAEDH8BBB7A363.

2.3.4 Capacitors

A digital circuit like an LED driver can change its current consumption significantly within a short amount of time. Because the connections to the power supply are relatively long, they have a substantial parasitic inductivity. So when the IC suddenly needs more current because all its LEDs have just been switched on, this current is not available instantly. Worse yet, if the LEDs are on and then are turned off, the current can't immediately stop flowing and will induce dangerous voltage spikes.

In order to prevent those two situations from happening, I used capacitors. Namely, each LED driver has its own 1 µF ceramic capacitor.

2.3.5 Motor

In order for the whole device to rotate, I needed a motor. Luckily, I could salvage an appropriate part from an old RC motor boat. It's a big and heavy brushed motor. The fact that it has a built-in fan and was water-cooled in the boat leads me to believe that it can take a lot more power than what I need for the display.

Because of its high torque, I could attach the motor directly to the rotating shaft instead of using gears or a belt drive. That way there are less moving parts and the motor can turn at a lower rotational speed, both resulting in less wear and noise.

2.3.6 Power Transmission

To transmit power from the stationary power supply to the rotating part, I used slip rings. I made a relatively simple version, consisting of two round pieces of copper sheet on the rotating end, and two long pieces of copper sheet on the stationary part.

The slip rings must be able to transmit the current necessary to light all LEDs up at once. Assuming that each LED needs 20 mA, this current is about 9.6 A. Additionally, there is a small current draw by the microcontroller, but it will be negligible in comparison. Testing has shown that the slip rings don't heat up noticeably when conducting 10 A for several minutes.

In practice, however, this current is not reached since the vast majority of pixels are switched off when displaying a typical 3D image.

Because two pieces of copper sheet rubbing against each other aren't the most reliable electrical connection, I used a 2200 µF electrolytic capacitor to smooth out the supply voltage. Even with it, I found that the pressure between the two copper conductors had to be quite high for them to form a good contact.

2.4 Circuit Diagram
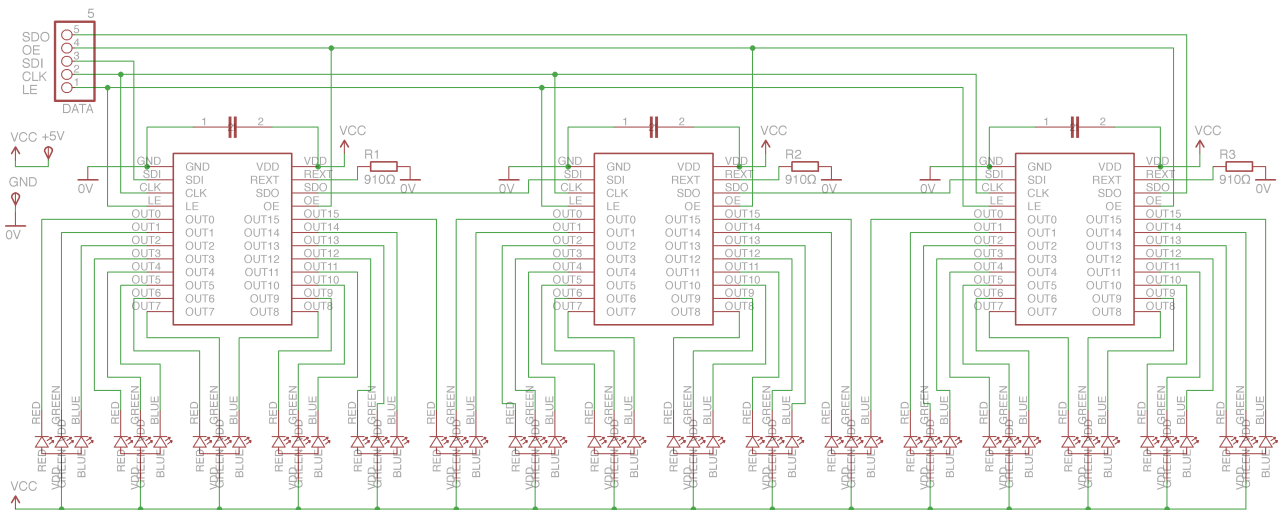
Figure 2: The finished circuit diagram

The circuit diagram — as well as the PCB layout — were created in Eagle [11]. Its free version is limited to two-layer boards of sizes up to 10 cm by 8 cm, but that's more than enough for my purposes.

In figure 2, the final version of the schematic is pictured. For the sake of simplicity and readibility, the layout corresponds roughly to how the board will look.

I had to configure the footprints for the LED drivers and the LEDs by myself, which I did by reading about the physical dimensions in the respective data sheet. For the resistors and capacitors, Eagle had pre-built footprints which I could use.

2.4.1 Data Interface

For transmitting data to the LED drivers, I used a simple row of 5 pins. Two of those pins are the data lines: The input is connected to the first chip's input, and the output comes from the last chip's output. The rest of the pins (clock, output enable and latch enable) are control lines, which go to all three chips simultaneously. The order of these connections was determined when making the PCB layout.

2.4.2 LED Drivers

The LED drivers are arranged in a so-called daisy chain configuration, where the serial data output of each chip is connected to the serial data input of the next chip. This is not only the case within one board, but also between the different boards: The data output of one board is connected to the next board's input. That way, all of the 30 LED drivers can be controlled by a single SPI bus without anything like a chip select signal. To send data to the chips, the microcontroller can simply send out 480 bit without interruption.

The external resistor which is connected to each LED driver sets the output current. The following equation, taken from page 15 in the datasheet, describes the relation between the output current (I_{OUT}) and the resistor's value (R_{EXT}):

$$I_{OUT} = 15 \cdot \frac{1.25\,\text{V}}{R_{EXT}}$$

Let's find out the resistance necessary for the output current to be $20\,\text{mA}$:

$$R_{EXT} = 15 \cdot \frac{1.25\,\text{V}}{20\,\text{mA}} = 937.5\,\Omega$$

The next smaller value within the E24 series is $910\,\Omega$. By choosing this resistance, the output current is $20.6\,\text{mA}$, which is close enough to the desired $20\,\text{mA}$. I ordered 40 through-hole resistors of this value, because at the time I had in mind to etch my own boards at home.

2.4.3 LEDs

The LEDs have their common anodes connected to V_{DD}. The cathodes are connected to a driver output each, and are arranged red - green - blue (from left to right). Series resistors aren't needed because the LED drivers handle current regulation.

2.5 PCB Design

2.5.1 General Remarks

The defining factor for the board layout was the pin layout of the LED drivers. The outputs are all at the bottom of the chip, while the data interface is at the top. This led me to the layout I have now: the LEDs are in a row at the bottom of the board where they can be directly connected to the driver outputs on the upper copper layer. The control lines are routed on the back side of the board behind the chips, with vias going to the front side to connect to each chip. The data lines go straight from one chip to the next one on the upper copper layer, with the last output going back to the board's data interface on the back side. The resistors are placed on the right side of their respective driver IC, jumping across the data traces.

The supply lines are placed beneath the row of LEDs. Had I placed the supply on the upper edge of the board, I still would've had to connect all the 16 anodes to V_{DD}. The driver chips, on the other hand, only need one connection to V_{DD} and GND each, so by placing the supply lines at the bottom of the board I have less traces running across the board vertically.
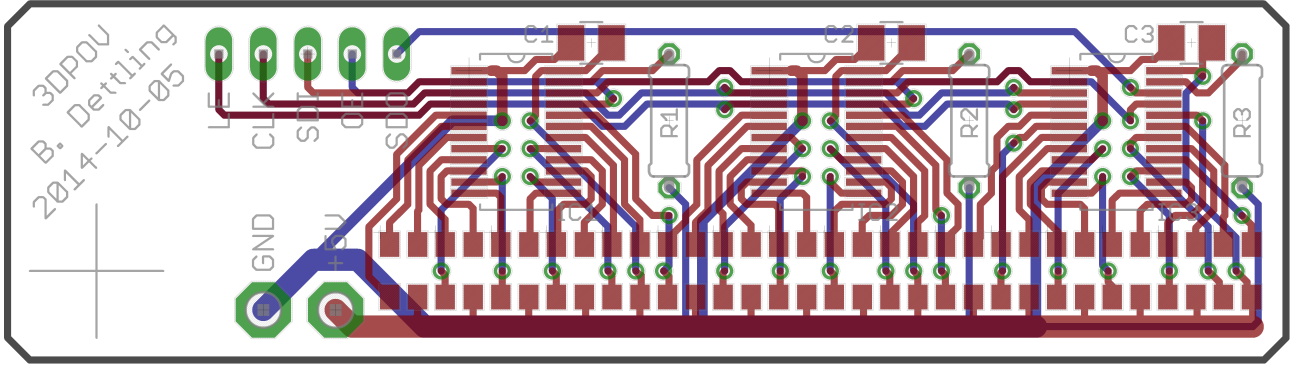
Figure 3: The PCB layout in Eagle. Red = upper copper layer, blue = lower copper layer, green = vias, grey = silkscreen.

From a layout perspective, it would've made sense to put the GND line at the top of the board, since the chips need one connection to that trace each, while the LEDs don't. However, I opted to put both supply lines right above each other to reduce electromagnetic interference. If all the LEDs are switched on suddenly and a lot of current starts to flow, a magnetic field builds up quickly around the supply traces. This could induce bothering or even dangerous currents in the surrounding electronics. By having the two supply lines close to each other, their magnetic fields largely cancel each other out.

2.5.2 Space-saving measures

In order for the production to be as cheap as possible as well as to minimise the rotating mass, I tried to make the boards as small as possible. Ultimately, I managed to bring the size down to $73\,\text{mm}$ by $21\,\text{mm}$.

In order to achieve this, but also to avoid large gaps between the pixels, I placed the LEDs as close as possible to each other. The datasheet claims a width of $2.8 \pm 0.2\,\text{mm}$, so I decided to place the LEDs at $3\,\text{mm}$ increments. On a grid of $12.5\,\text{mil}$, the next larger distance is $125\,\text{mil}$, which is equal to $3.175\,\text{mm}$.

The connections to the cathodes of the red LED needs to pass by the other two cathodes, since they're located on the lower side of the LED package. Because the LEDs are so close to each other, the only option was to use the lower copper layer. The only place to go back to the upper copper layer was right beneath the driver chips.

Initially, I had placed the decoupling capacitors right above the ICs. However, I could save a small bit of vertical board size by placing them at an offset to the right. That way I could push them down until where the pins begin, which is a bit further down than the edge of the chip itself.

2.5.3 Thermal Considerations

A positive side effect of placing some vias beneath the chip is their thermal conductivity, which helps disspiate the chip's excess heat. The maximum power output of the chips — when all the attached LEDs are turned on — is as follows:

$$\text{P}_{\text{MAX}} = (5\,\text{V} - \text{V}_{\text{F}}) \cdot 20.6\,\text{mA/output} \cdot 16\,\text{outputs} = 780\,\text{mW}$$

V_{F} is different depending on the colour of the LED, so I just used the average value, which is $2.633\,\text{V}$. The actual values are $1.9\,\text{V}$ for red and $3\,\text{V}$ for green and blue.

A power output of $780\,\text{mW}$ is not to be underestimated on such a small PCB, and tests have shown that that the boards do get quite hot when operated at full power. In reality, however, the situation is much less severe because the LEDs are typically turned on only during a fraction of the rotation. Furthermore, the quick rotation results in excellent air circulation, which aids the dissipation of heat.

3. Fabrication

3.1 Circuit Boards

3.1.1 Fabrication

There's no shortage of firms who produce PCBs according to a custom layout. One of the cheapest providers to ship small production runs within a reasonable timeframe and with good production quality turned out to be OSHPark [12]. They only produce in multiples of 3, so I ordered 12 instead of 10 (which I'd have done anyway to have some spare boards). I submitted my order on the 5th of October 2014 and paid US$ 47. Ten days later I received a message that the boards were finished, and after another ten days of waiting they arrived at my doorstep.


Figure 4: A finished circuit board

3.1.2 Assembly

As soon as I got the boards in the mail, I started soldering the parts to them. I started with the LEDs to avoid having to solder in the tight gap that would later be formed between the LEDs and their driver chips. It took me two hours to solder the first board, but with increasing practice the time for one row of LEDs got down to 25 or 30 minutes.

The chips were a lot easier to solder than the LEDs, despite the distance between the individual pins only being 0.635 mm. I could simply drown all the pins in solder, and then wick away all the solder bridges. That way I could solder twelve pins within the blink of an eye.

3.1.3 Testing and Debugging

With the boards finished, I wanted to know whether or not they work. In order to find out, I wrote a little program on the microcontroller to light the LEDs up in all possible colours consecutively. It did not work at the first try, so I used an oscilloscope to see if data was being sent at all. This was the case, the SPI library on the Teensy performed its task flawlessly.

It turned out that the culprit were the bad contacts. Instead of soldering the wires onto the board and the microcontroller, I simply inserted them into the respective holes. This caused the connection

to have too much parasitic capacitance, and the signals (which were being sent at 4 MHz) did not have time to propagate up to their destination.

I soldered the cables on properly and sure enough, the board worked as intended. I repeated this procedure for all ten boards to see if they all worked as well as the first one, which wasn't the case.

Because of reasons unbeknownst to me, one particular LED driver started turning on all of its LEDs at a moderate brightness as soon as I powered it up — even though no data was being sent at all. At first I thought that the serial data input had a solder bridge pulling it up to V_{DD}, and immediately sending a whole lot of 1's upon powering up the board. However, the data input pin was not close enough to a positive supply line for this to be an issue. Additionally, there would have to be a clock signal as well as a pulse on the latch pin for the LEDs to light up. I gave up and assumed that the chip was faulty, and when I replaced it with a spare part everything worked perfectly.

Apart from that, there were no major issues. I had a few LEDs which didn't light up when they should have. One of them had a solder bridge between the common anode and the cathode of the red LED, which made a voltage drop across the red LED impossible. Two were simply soldered on poorly and weren't connected to the drivers at all. A few other LEDs didn't work and I found no obvious issues. I replaced them with spares, which did the trick.

3.2 Physical Assembly

3.2.1 Frame

The frame was built out of beech wood strips with a square cross section and a side length of 1.5 cm. The baseplate is made out of plywood and has the dimensions 15 cm by 15 cm by 1 cm. On it, I mounted two vertical wood strips with a length of 20 cm, which were both stabilise by three struts. A horizontal board, mounted to the vertical strips at a height of 20 cm, holds the motor.

On the top of the two vertical strips I placed another strip to hold the shaft in its position. A simple hole served as a bearing, which I later replaced with a plastic insert (kindly sponsored by IGUS).

3.2.2 Slip Rings

In order to supply the rotating parts with electricity, I mounted a wooden disc right at the bottom of the shaft. On both sides of that disc I glued on a round piece of copper sheet, which I then connected to cables going up to where the electronics would later be.

On the stationary end, two long pieces of copper sheet serve as contacts. Because of their spring-like properties, they exert a certain amount of pressure on the rotating disc.

Despite this, a considerable amount of pressure was necessary to form a good contact. The contact pressure proved to be a primitive, but usable way of roughly adjusting the rotational speed of the display. Sadly it's quite an inefficient approach: The motor uses a lot of current when friction is applied to its load. It becomes noticeably warm after a few minutes of use, and measurements on the current limiting resistors have shown that it draws about 6.5 A at 5 V.

This constant pressure, in addition to the fact that copper is quite a soft metal, results in considerable wear on the slip rings. Luckily, the stationary ends can be replaced without much hassle. The rotating part would be a pain to replace, but since wear is distributed across the whole copper disc, it will be a long time before this needs to be done.

3.2.3 Controller

The microcontroller is mounted on another wooden disc, located a few centimetres above the slip rings. It also hosts the hall sensor, which is used to measure the rotational speed and calculate the

position of the rotating part at any given moment. The corresponding magnet is mounted on a piece of threaded rod, which itself is mounted on the same board as the motor.

Additionally, the 2200 µF capacitor is located on the upper wooden disc. In the end, this place has become a bit of a mess, especially with all the cables going up to the circuit boards. Luckily this isn't visible while the apparatus is running.

3.2.4   Circuit Boards

The boards were already equipped with components, but what was missing was the hold to mount them on the shaft. The diameter of this shaft is 4 mm. In order to find a good hole size, I drilled a few test holes on a spare board. A hole of 4 mm was so small for the boards to fit on the shaft, while 4.5 mm was obviously too big. After chamfering the top end of the shaft, the board with a 4 mm hole just barely could slide on the shaft. This tight fit proved to be advantageous, because torque can now be transmitted reliably from the shaft to the boards, and they don't slide around.

There were two factors I had to carefully consider while arranging the boards:

- The boards had to be balanced out to minimise static and dynamic unbalance

- The amount of boards obstructing other boards had to be minimal

In the end, I decided to arrange the boards in a double-spiral arrangement. This was the only arrangement I could think of that (theoretically) has no static imbalance at all. Because each board has another board on the exact opposite side of the shaft, there are five pairs of boards where one board statically balances out the other. However, each pair has some dynamic inbalance because the two boards aren't mounted at the same height. This could've been minimised by alternating the order of boards between each pair. However, this wasn't necessary because the vibrations were within reasonable bounds.

The visibility is not too bad either: On the viewer-facing half of the display, all LEDs are constantly visible. On the backside, it is possible that one group of boards is right before the other one, causing some layers of the image to disappear.

4. Programming

4.1 Used Software

The following program has been compiled with the Arduino IDE 1.0.6 [13] and sent to the Teensy by the Teensyduino 1.20 plugin [14]. However, the most part of the program has been written in Sublime Text 2 [15] with the stino plugin [16]. This combination has a few advantages over simply using the Arduino IDE to write all code, namely there's automatic code completion, support for several cursors, and a search-and-replace function with regex suport. I installed and used these programs on Mac OS 10.9 and later 10.10.

4.2 Timing

4.2.1 Theory

In order for the image to stay still despite the rotational speed not being completely constant, the frequency of LED refreshed must be constantly adjusted based on the rotational period. Because there are 100 pixels per revolution, the time to display one pixel must be $\frac{1}{100}$ of the rotational period.

This value can be measured easily by using the hall sensor. By dividing this value by 100, we can obtain the duration of one pixel for the current speed. If the rotational speed changes, that's a bit of a problem: The measured value is from the previous rotation, but in the current rotation, the device already turns slightly faster. This causes the image to be slightly warped right after powering up the display. However, once it has reached temrminal velocity, this effect is no longer noticeable, because now the subsequent rotations have about the same period.

4.2.2 Implementation

This behaviour is implemented with an interrupt that runs each time the magnet is detected (i.e. after each rotation). Within that interrupt, several things happen:

First, we read a variable called `sinceMagnet`. As its name says, it keeps track of the number of microseconds since the last time the magnet was detected. It's of the special type `elapsedMicros`, which means that it's automatically incremented each microsecond. This value is divided by 100 and stored in `microsPerPixel`, where it represents the amount of time that any given pixel should be displayed for.

This variable is now reset to 0 to start the new rotation. The same is done to another variable, `nextPixelMicros`. It contains the next point in time when the pixels should be refreshed, with respect to the start of the rotation (which just happened).

In the main loop of the program, this variable is continuously compared to `sinceMagnet`. If `sinceMagnet` is bigger, the LEDs are refreshed by calling `sendData`. Within that function, `nextPixelMicros` is incremented by `microsPerPixel`, causing the program to wait for the display to turn another $\frac{1}{100}$ of a full circle. After that, the next pixel will be displayed.

The implementation of the described behaviour can be found in the appendix, where the whole source

code is included. Specifically, what I've just talked about is contained in the `timerUpdate` function, as well as in the main `loop`.

4.3 Data storage and display

The previously mentioned function `sendData` needs to be able to know which bytes to send based on one parameter, which is the current angle of of the display. In this section, I will explain how I achieved this.

To save the image on the microcontroller, I used a simple three-dimensional array containing $100 \cdot 10 \cdot 6$ bytes of data. The first index stands for the angle of a pixel in hundredths of a circle, the second index corresponds to the height and the third one describes the position of each byte on a row of LEDs. Because there are 48 LEDs per row, 6 bytes can be used to describe the row's state without wasting any storage space.

Because the PCBs are not right above each other, the data that is being sent out at the same time don't correspond to the pixels that are above each other in the image. In order avoid having to account for the angular offsets of each board while creating an image, I made the program do that work automatically.

Each pair of two boards are mounted on the exact opposite side, making the angle between them 180° or 50 pixels. Each subsequent pair is additionally rotated by 28.8° which is equivalent to 8 pixels.

The program accounts for this arrangement by using a simple for loop. It iterates over five pairs of boards. At its start, the program calculates the angle and height of the first board based on the control variable. After the data has been sent, it then moves to the second board by adding 50 to the angle and incrementing the board counter, and feeds the second board with data too.

After these two additions, the angle value will go over 99 in many cases, which is the highest possible array index. Since the pixel at position $100 + n$ is the same as the one at position n, I can pervent causing undefined behaviour by applying a modulus 100 operator to the angle value before using it to access the array.

4.4 Miscellaneous

While the program works the way it is described above, there are a few small things that I managed to further improve.

Because I selected the position of the boards with respect to the hall sensor arbitrarily, the image was turned around. To fix that, I added a constant value to all angle values, so now the first pixel is displayed exactly at 0 instead of whereever the hall sensor detects the magnet.

Whenever I started up the program, it started displaying the first pixel. It even did so when I connected the Teensy to my laptop via USB. This was dangerous because neither the Teensy nor the laptop's USB port can handle a lot of current, and if more than a handful of LEDs were switched on, the Teensy could've been damaged.

To fix this I added a boolean variable called `running`, which is set to `false` by default and only flipped when the magnet is first detected. That way the program only starts running if the display is actually rotating, at which point we can be sure that power is supplied through the slip rings rather than the USB port.

The images displayed were constantly jittering around instead of staying still like they should. It took me a long time to find the source of this bug, but I managed to fix it in time for the swiss national competition organised by "Swiss Youth in Science".

The problem was that sometimes, the magnet was detected while `sendData` was still running. Because the hall sensor's interrupt had a higher priority than the timer interrupt I had used to send data, control was immediately given to the former. The `timerUpdate` function was executed normally,

resetting `currentPixel` to zero and reading out the duration of the last rotation. However, after that, control was given bach to the `sendData` function. At the end of that function, `currentPixel` was incremented by one. That way, the `currentPixel` variable already had a value of one when it should've stayed at 0. For the entirety of the next rotation, `currentPixel` was off by one, resulting in the image quickly turning by 3.6°.

4.5   Image Creation

Initially, all my images were made completely by hand. I had to enter the individual bits into the program that would result in the desired image. Obviously this was a time-consuming and error prone task. In order to simplify the process of making an image, I wrote a python program that would support me in doing so. All of the source code is available on `https://github.com/mbjd/_3DPOV` in the folder `image_creation`.

4.5.1   Basics

In order to save an image, I used a three-dimensional array of $100 \cdot 10 \cdot 16$ integers. The lowest three bits of each integer represent the states of the red, green and blue LED respectively. The rest of the bits are unused. This causes the program to use more memory than strictly necessary, but it makes editing the pixels much simpler.

On the display itself everything stays the same, and no bits are wasted. The conversion from the easily editable to the memory efficient format is automatically handled by the python program.

The pixels are arranged in a polar coordinate system. Because it's simpler to express straight lines, spheres or surfaces in terms of cartesian coordinates, calculations typically go like this:

forall the *pixels in the image* **do**
  convert polar coordinates to cartesian;
  calculate some function with the coordinates as input;
  colour the pixel accordingly;
end

4.5.2   Straight lines

Given two points, the program can draw a line connecting them. To do that, it calculates the distance to the line for each point, and if it's below a certain threshold, the point is coloured. This threshold controls the thickness of the line and can be selected freely, as can the colour of the line. The implementation can be found in the functions `drawLine`.

4.5.3   Spheres

Spheres are even simpler than straight lines: For each point, the program compares the distance to the center against the desired radius of the sphere. If it's smaller than the radius, the point is within the sphere and can be coloured. This functionality is implemented as `drawSphere`. Another version of this function, called `hollowSphere`, checks if the distance to the middle point is within a certain range, drawing a hollow sphere.

4.5.4   Planes

A plane can be defined in terms of the equation $ax + by + cz = d$, where x, y, and z are the cartesian coordinates, and a, b, c, and d are parameters defining the surface. Additionally, the function takes an argument specifying the thickness of the plane. In the end, all pixels whose coordinates (x, y, z)

satisfy $|ax + by + cz - d| < \frac{\text{thickness}}{2}$ are coloured with the specified colour. The implementation is called `drawSurface`.

4.5.5 Functions

There are two fundamental types of functions that the program can draw. The first one are functions of the form $f(x, y, z) \rightarrow colour$. In order to render such a function, I simply iterate over all pixels in the image, calculating the function and colouring the pixel with its return value.

Secondly, the program can plot functions of the form $f(x, y) \rightarrow z$. It is a special case of the first one and could be defined in terms of it: I could simply iterate over all 16000 pixels and check if $\lfloor f(x, y) \rfloor = z$. However, that would be inefficient, because the same calculation would be made 10 times for the 10 different z values. Instead, I only iterate over all combinations of x and y values, and colour the pixel at the corresponding z value. That way, the function is only evaluated 1600 instead of 16000 times.

The respective implementations can be found in `plotColourFunction` and `plotRealFunction`.

There's a third one called `plotBoolFunction`. It's similar to `plotColourFunction` and can also be redefined in terms of it (with no significant loss in performance). Instead of converting the coordinates (x, y, z) to a colour, it takes a function that returns a boolean value, and a colour. If the result of the given function is `True`, the pixel is coloured with the given colour, otherwise nothing happens.

4.5.6 Additional functionality

Using these basic building blocks, I can draw more complex shapes by repeatedly calling one of the above functions with different arguments. For example, there's the function `connectcircle`. It takes a list of points and connects each with the next one, and the last to the first point. That way I can draw arbitrary polygons.

I can also automatically draw a cuboid by specifying two points. The cuboid is then put up between these points, all edges being parallel to the coordinate axes. The implementation is called `drawCuboid`.

Another function, `connectAll`, takes a group of points and draws a line between each possible pair of points. That way, I can draw a 3-sided pyramid in one go.

4.5.7 Usage

This whole program is a python module, which is to say you can import it from any other python program, including the interactive python shell.

```
$ python
[...]
>>> from mkmodel import *
```

After doing that, all of the functions defined in the program are available to play with. Using the `newImage` function, an empty image can be generated. We then draw a line and a sphere in the image:

```
>>> img = newImage()
>>> drawLinePolar(img, (0, 7, 14), (30, 2, 14), green, 0.5)
>>> drawSpherePolar(img, (62, 4.5, 7), blue, 4.5)
```

When the desired image is created, it can be saved as an Arduino program:

```
>>> writeSketch(img, 'example')
```

The complete program, ready to upload to the Teensy, is then saved in `./example/example.ino`. Creating a directory for each program is necessary because the Arduino IDE only opens them if they're located in a directory with the same name as the program.

5. Operation

5.1 Power Supply

Initially I had planned to supply the PCBs with 5 V. However, the Teensy sends its signals at 3.3 V. The driver chips, on the other hand, only read a logical high when a signal is over $0.7 \cdot V_{DD}$. At 5 V, this value would be 3.5 V. So I tried powering the boards with only 3.3 V, and everything worked perfectly. Better yet, the boards weren't getting as hot anymore.

According to the manufacturer, the Teensy needs at least 3.7 V to operate, presumably because it uses a linear regulator to get a stable voltage ov 3.3 V. However, there were no problems when I started using 3.3 V. Even though the regulator probably reduces that voltage to less than 3 V, the processor still runs at its full clock speed of 96 MHz.

Unfortunately, the PC power supply which I used for the project broke down. As I later found out, the culprit was the motor consuming too much current. At that time, I didn't use any series resistance to limit the current flowing through the motor, which I powered with 3.3 V. This has caused the inrush current to be greater than what the power supply could handle.

To remedy this, I used a pair of long, thin wires to connect to the motor to provide some resistance. I later replaced these with two 150 mΩ power resistors. Now there are two ways to power the motor: I can use 5 V and connect the resistors in series for a total resistance of 300 mΩ, or I can power it with 3.3 V and connect the resistors in parallel, resulting in a resistance of 75 mΩ. The latter method is more efficient, but it causes the acceleration at startup to be slower.

5.2 Mechanical Problems

Because the slip rings weren't fixed in their position, they could slide out of their slots after a few minutes of operation. I fixed this issue by mounting a small wooden board onto these slots (perventing lateral movement), and routing the cables going from the power supply to the slip rings through two holes in that board (perventing lengthwise movement).

The boards weren't exactly aligned at the 28.8° that I specified. This caused the image to be displayed crookedly: Pixels that should've been right above each other were shifted sideways by a small distance. Whenever I adjusted the boards, they would be misaligned again after using the display for some time. To fix this, I aligned them correctly one last time and used enough solder on the power cables running through the boards for them to form a solid connection.

The bolt nut that held the motor in place started coming loose because of the vibrations. To pervent that from happening again, I used a second nut and tightened it against the first one.

6. Conclusions

6.1 Possible Improvements

6.1.1 Controller and Data Transmission

The Teensy 3.1 is currently mounted quite loosely on its wooden disc. A piece of wire goes through two unoccupied contacts as well as through two holes in the disc. In addition, the cables forming the SPI bus aren't laid out neatly at all; many are too long and don't take the direct route.

This situation could be improved by replacing the upper wooden disc with another printed circuit board. All connections could be routed elegantly on the PCB, and the SPI interfaces could be formed by two ribbon cables. The hall sensor could be properly soldered on instead of fixing it in place using hot glue.

Between the individual LED boards, ribbon cables could also be considered. However, the only connections that always go to the board right above are CLK, LE and OE, while the serial data inputs and outputs have to be connected to the board on the opposite side. These connections would still have to be formed by individual cables.

6.1.2 Slip Rings

Because the slip rings are quite rough and not completely horizontal, a lot of pressure is needed to form a consistent contact. This leads to high friction between the two contacts, which then results in quick wear as well as a lot of friction.

The main reason why the display is not energy efficient at all are the slip rings. At the start of the project, I had the idea of using a 3.5 mm phone connector to transmit power. However, I abolished that plan because of two reasons. Firstly, I had concerns that such a small connector would have trouble conducting the maximum current of 10 A. This would most likely not have been a problem, since only a small part of all pixels are turned on in a typical image, which causes the average current to be a small fraction of the maximum value.

The second problem would've been the difficulty of mounting a phone connector on the display. I would've had to place it at the end of the shaft, where it would've also served as a bearing. In order to minimise the forces acting upon the plug, it would've had to be mounted exactly in the center of the shaft. I couldn't come up with a solid way to connect both of the phone plugs' ends to the display.

A much more elegant solution would've been to use induction to transfer power. By having a stationary and a rotating coil close together, I could've passed a high-frequency alternating current through the stationary one, causing current to flow in the rotating one as well. The advantages of such an approach would be substantial: The bearings would be the only contact point between the rotating and stationary part, causing friction to be minimal. I could've also completely eliminated the annoying, broadband noise currently emitted by the slip rings.

I didn't implement wireless power transfer because I felt like the extra work wasn't worth it, especially considering this was a prototype. If I wanted to produce this display en masse, inductive power transfer would definitely be the way to go as it's the only way to make the display somewhat quiet.

6.1.3 Motor

I decided to drive the display directly without any kind of gear reduction because that was the simplest solution. I only could do so because the motor was strong enough. Unfortunately, it consumes a substantial amount of current.

If I wanted to sell my display commercially, I'd use a smaller motor, drive it at a higher rotational speed and gear it down. To save energy, I would also implement one of the above imporvements concerning power transmission.

6.2 Future Outlook

6.2.1 Animated images

Something I had in mind since the beginning of this project were animated images. The Teensy 3.1 has 256 kB of memory, which would fit 42 images of 6 kB each. This corresponts to roughly a second of playback. In order to display longer animations, there are several workarounds: Instead of showing one image that spans the entirety of the rotation, I can instead only save an image that is 25 pixels and display it four times.

In order to get around the limit of flash memory, external storage would be necessary. Even with only 1 GB, I could save 166666 frames, thus removing any practical limit in terms of playback time.

6.2.2 Interactive Programs

It would be interesting to not only display predefined images, but to actually run a program that can respond to user input, like a video game. In order to communicate with the external world, a wireless interface would be necessary. Readily available solutions include small WLAN or bluetooth modules, or the transmitters and receivers commonly used in radio-controlled aircraft.

6.2.3 Applications

Because of the low resolution and colour depth, the high level of noise, quick wear and the inability to send images to the display while running, this display is not suited as a replacement for your everyday computer or TV screen. However, these problems can all be fixed completely or partially, so that there could one day exist a version of my display which is actually suitable for the above purposes.

One domain where I could imagine a display like mine being especially useful is advertising. Because most people haven't ever seen such a display before, it would be an immediate eye-catcher in the display window of any business.

Another fun thing to try would be video games. Although the limited resolution and computing power would make it difficult to run the latest and greatest games, 3D versions of classics like Pong, Snake or a simple platform game would certainly be worth a shot.

6.3 Theoretical Limits

The current version of the display works fine and well, but it's certainly not at the limit of what can be done. In this section, I try to find out which aspects of the display can be improved how much.

6.3.1 Resolution

The resolution was determined based upon the size of the LEDs I used. Pixels are located 3.125 mm apart from each other, while the LEDs themselves have a width of about 2.8 mm. In order to decrease

pixel size significantly, I'd have to use smaller pixels. To a certain extent, this would be possible; there are LEDs that are about half the width of the ones I use currently.

With such small LEDs, there wouln't be enough space for the LED drivers. I didn't find chips with similar functionality in a smaller package. To a certain extent it would be possible to minimise the necessary space between LED drivers by using a multi-layer board layout. This method might just be enough for twice the pixel density, but more won't be possible. Beyond that, custom-made LED drivers are the only solution I can think of.

There's one more approach to increasing pixel density, which is to just have one colour per pixel.

Bibliography

[1] *POV-Display*, project report, retrieved 2016-05-19,
    `http://www.mikrocontroller.net/articles/POV-Display`

[2] *ATtiny85 POV Display*, build instructions by vishalapr, retrieved 2016-05-19
    `http://www.instructables.com/id/ATtiny8545-POV-Display/`

[3] *Build an LED Propeller Clock*, build instructions by Bob Blick, retrieved 2016-05-19
    `http://www.jameco.com/Jameco/workshop/MyStory/diy-led-propeller-clock.html`

[4] *voLumen - volumetric 3D display device*, project report by Maximilian Mali, retrieved 2016-05-19
    `https://hackaday.io/project/1737`

[5] *3D POV Display... Work in Progress*, video by Jason Hotchkiss, retrieved 2016-05-19
    `https://www.youtube.com/watch?v=rlptJzWwCXg`

[6] *Persistence of Vision* on Wikipedia, retrieved 2016-05-19
    `https://en.wikipedia.org/wiki/Persistence_of_vision`

[7] *Teensy USB Development Board*, product page, retrieved 2016-05-19
    `https://www.pjrc.com/teensy/`

[8] *New I^2C Library for Teensy 3*, forum post by nox771, retrieved 2016-05-19
    `https://forum.pjrc.com/threads/21680`

[9] *TLC5927*, product page, retrieved 2016-05-19
    `http://www.ti.com/product/TLC5927`

[10] *Texas-Instruments TLC5927IDBQR*, retrieved 2016-05-19
    `http://www.digikey.ch/product-detail/en/TLC5927IDBQR/296-24762-1-ND/`

[11] *What is EAGLE*, product page, retrieved 2016-05-19
    `http://www.cadsoftusa.com/eagle-pcb-design-software/about-eagle/`

[12] *OSH Park — An electric ecosystem*, retrieved 2016-05-19
    `https://oshpark.com/`

[13] *Download the Arduino Software*, retrieved 2016-05-19
    `https://www.arduino.cc/en/main/software`

[14] *Download Teensyduino*, retrieved 2016-05-19
    `https://www.pjrc.com/teensy/td_download.html`

[15] *Sublime Text 2*, product page, retrieved 2016-05-19
    `http://www.sublimetext.com/2`

[16] *Stino*, version control page by Robot-Will, retrieved 2016-05-19
    `https://github.com/Robot-Will/Stino`

[17] *Using Digital I/O Pins*, documentation page, retrieved 2016-05-19
https://www.pjrc.com/teensy/td_digital.html

[18] *SPI Library*, documentation page, retrieved 2016-05-19
https://www.pjrc.com/teensy/td_libs_SPI.html

[19] *Teensy 3.1 Pins*, retrieved 2016-05-19
https://www.pjrc.com/teensy/pinout.html

[20] *Phone Connector (audio)* on Wikipedia, retrieved 2016-05-19
https://en.wikipedia.org/wiki/Phone_connector